

Outclass Proposal for team Automode

Team name: Automode (Shouryamaan Jain & Arag Agrawal)

Product Name: Cohora

Problem Statement & Target Users

Campus community members can't quickly discover who on campus has specific skills or experience because expertise is scattered across unstructured groups, projects, and chats.

Cohora maps authenticated campus activity into a searchable, explainable knowledge graph that returns precise people-searches and shortest-path introductions.

The target users are students, faculty, and staff who need quick, explainable ways to find campus expertise and the shortest connection paths (*see chart*).

Segment	Description	Motivation
Students	Across cohorts and tracks	Want to find peers for hackathons, startups, or mentorship
Faculty	Professors and course mentors	Need to identify students with specific technical or analytical strengths
Staff/Clubs	Admins and student bodies	Recruit contributors for initiatives or projects

Scope of project

In-scope: Build a consented campus knowledge graph from profiles, posts, projects and group memberships; answer natural-language queries with top-K people, provenance, and shortest-path visualizations; iteratively improve rankings from user feedback using vector DB + memory.

Out-of-scope: Drafting or sending outreach/automated introductions; full hiring workflows or formal background verification; profiling or surfacing off-campus data without explicit user consent.

Typical use case

Student project — A student needs teammates with a specific skill (e.g., [authentication](#)). They type a plain-English request; Cohora returns a ranked top list with one-line evidence for each person ([project](#) / [post](#) / [repo](#)) and the shortest intro path to each candidate.

PR team — A PR lead needs students with national-level sports credentials. They ask Cohora, which returns verified matches with visible evidence (certificates, club posts), shows relevant group memberships, and the shortest consent-respecting contact paths.

Shared goals — Fast answers, clear evidence for every match, and actionable, privacy-respecting connection paths.

How it works — Cohora reads the query with an LLM, finds similar profiles via semantic search, uses the graph to compute shortest paths, and learns from user feedback to improve results.

insights from user interview

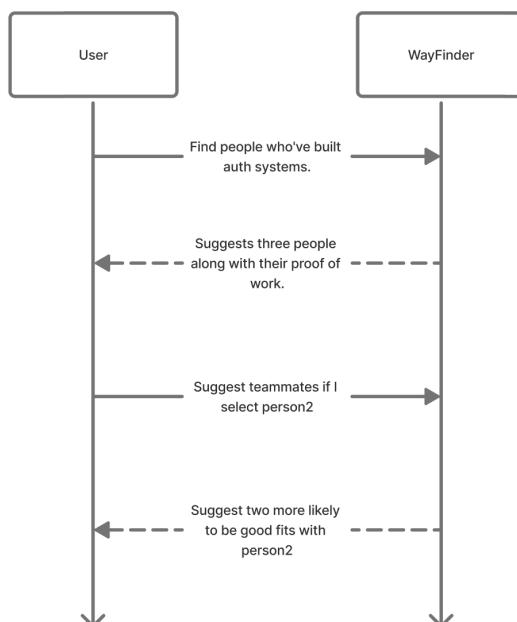
1. Many users don't know who on campus has specific technical or domain skills beyond their immediate cohort.
 2. People rely on word-of-mouth and scattered posts, so discovery is slow and often misses good matches.
 3. Users want quick, evidence-backed results (*a short proof or post*) rather than vague suggestions.
 4. Privacy and consent matter—users will share data if they can control visibility and opt out of specific uses.
 5. A simple chat interface with a visual connection map and clear intro paths is the most usable format for discovery.
-

About chatbot

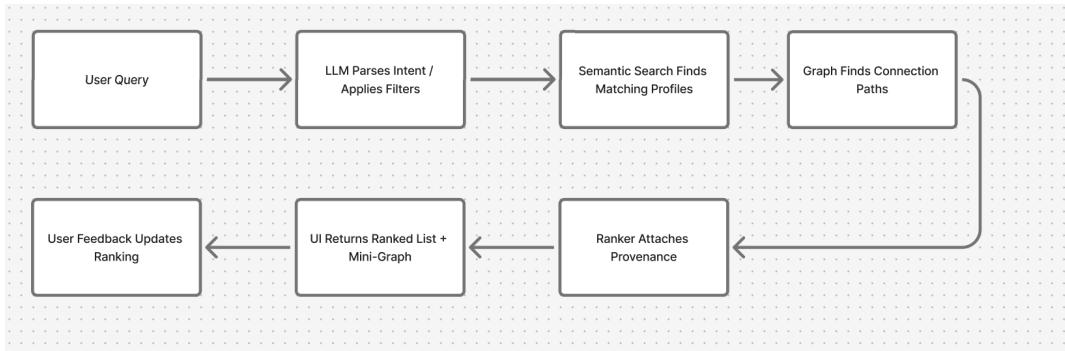
Persona

Concise, curious, and evidence-first, the chatBot speaks like a socially active friend: tailored to needs, direct and always shows *why* it recommends someone. This persona fits because users want fast, practical answers with verifiability.

Sample dialogue



Flow diagram



Core feature and functionality

Feature	Conversational discovery chatbot	Verifiability & evidence layer
Does	Accepts plain-English queries, returns top-K people with one-line proof and shortest contact path; supports simple filters/follow-ups.	Attaches concrete evidence to each match (repos, posts, certificates, groups) with a confidence tag; users can flag/confirm.
Value	Finds credible collaborators or experts across campus in seconds.	Provides trust and quick validation so results are immediately actionable.
Uses / Integrations	LLM for intent parsing; vector DB for semantic matching; graph DB for pathfinding; OAuth to campus directory/GitHub/LinkedIn.	Data connectors (GitHub/LinkedIn/Twitter, campus data); parsers/OCR; LLM for summarizing evidence.

Ethical guardrails and success metrics

We will protect user privacy and avoid doxxing or surfacing sensitive attributes. Cohora will label recommendations as inferences and only show private profiles, connection chains, or documents with explicit consent. The main guardrail is a consent-first visibility toggle ([public](#) / [campus-only](#) / [hidden](#)). And a short disclaimer on every result.

We'll track a few simple metrics to judge progress — for example, the share of top-5 suggestions rated relevant (aim ~70%), median response time (target under <10s), average satisfaction (around 4/5). These are starting targets to monitor and refine.